

# Java Threads (cont'd)

---





# Agenda

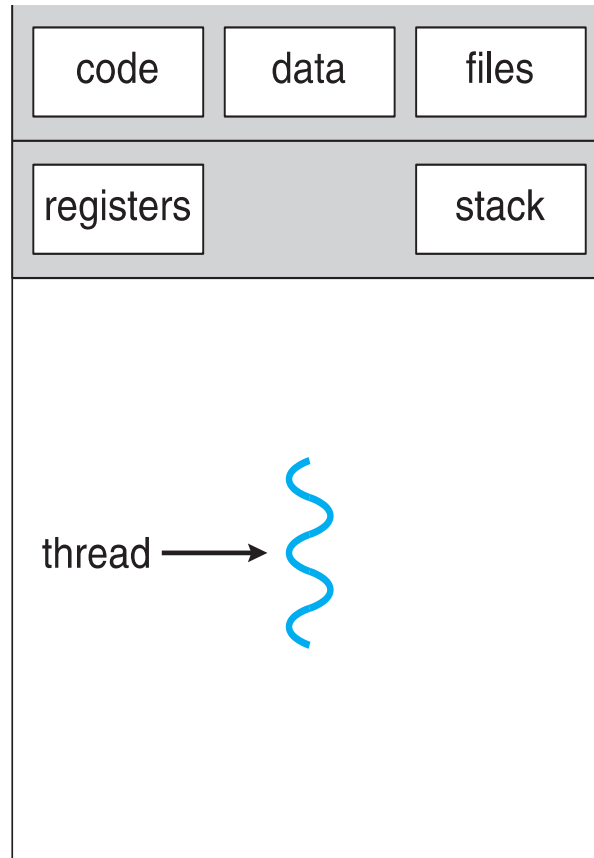
---

- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading

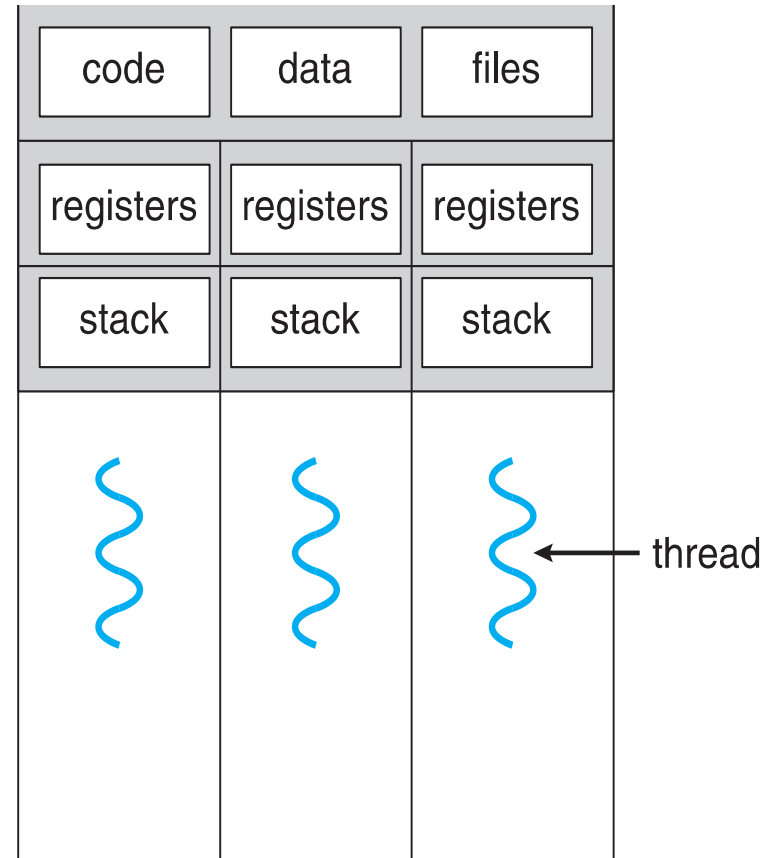




# Single and Multithreaded Processes



single-threaded process



multithreaded process





# Benefits

---

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multiprocessor architectures





# Multicore Programming

---


- **Multi-CPU systems.** Multiple CPUs are placed in the computer to provide more computing performance.
- **Multicore systems.** Multiple computing cores are placed on a single processing chip where each core appears as a separate CPU to the operating system
- Multithreaded programming provides a mechanism for more efficient use of these multiple computing cores and improved concurrency.
- Consider an application with four threads.
  - On a system with a single computing core, concurrency means that the execution of the threads will be interleaved over time.
  - On a system with multiple cores, however, concurrency means that some threads can run in parallel, because the system can assign a separate thread to each core






# Concurrency vs. Parallelism

## ■ Concurrent execution on single-core system:



## ■ Parallelism on a multi-core system:





# Multicore Programming (Cont.)

---

- There is a fine but clear distinction between concurrency and parallelism..
- A concurrent system supports more than one task by allowing all the tasks to make progress.
- In contrast, a system is parallel if it can perform more than one task simultaneously.
- Thus, it is possible to have concurrency without parallelism





# Multicore Programming (Cont.)



- Types of parallelism
  - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
  - **Task parallelism** – distributing threads across cores, each thread performing unique operation
- As number of threads grows, so does architectural support for threading
  - CPUs have cores as well as ***hardware threads***
  - Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core







# Data and Task Parallelism





# Multicore Programming

---


- **Multicore** or **multiprocessor** systems are placing pressure on programmers. Challenges include:
  - **Dividing activities**
  - **Balance**
  - **Data splitting**
  - **Data dependency**
  - **Testing and debugging**
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
  - Single processor / core, scheduler providing concurrency





# User and Kernel Threads

- Support for threads may be provided at two different levels:
  - **User threads** - are supported above the kernel and are managed without kernel support, primarily by user-level threads library.
  - **Kernel threads** - are supported by and managed directly by the operating system.



- Virtually all contemporary systems support kernel threads:
  - Windows, Linux, and Mac OS X





# Relationship between user and Kernel threads

---


- Three common ways of establishing relationship between user and kernel threads:
  - Many-to-One
  - One-to-One
  - Many-to-Many





# One-to-One Model


- Each user-level thread maps to a single kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows
  - Linux





# Many-to-One Model


- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
  - **Solaris Green Threads**
  - **GNU Portable Threads**





# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package





# Thread Libraries

---

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS
- Three primary thread libraries:
  - POSIX **Pthreads**
  - Windows threads
  - Java threads







# Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Two techniques for creating threads in a Java program.
  - One approach is to create a new class that is derived from the Thread class and to override its run() method.
  - An alternative is to define a class that implements the Runnable interface, as shown below

```
public interface Runnable
{
    public abstract void run();
}
```

- When a class implements Runnable, it must define a run() method. The code implementing the run() method is what runs as a separate thread.





# Implicit Threading


---

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Methods explored
  - Thread Pools
  - Fork Join
  - OpenMP
  - Intel Thread Building Blocks
  - Grand Central Dispatch
- Other methods include Microsoft Threading Building Blocks (TBB), **java.util.concurrent** package






# Fork-Join Parallelism





# Fork-Join in JAVA





# Fork-Join Calculation using the Java API

---

```
import java.util.concurrent.*;
public class SumTask extends RecursiveTask<Integer>
{
    static final int THRESHOLD = 1000;
    private int begin;
    private int end;
    private int[] array;
    public SumTask(int begin, int end, int[] array) {
        this.begin = begin;
        this.end = end;
        this.array = array;
    }
}
```





## Fork-Join Calculation using the Java API (Cont.)

```
protected Integer compute() {  
    if (end - begin < THRESHOLD) {  
        int sum = 0;  
        for (int i = begin; i <= end; i++)  
            sum += array[i];  
        return sum;  
    }  
    else {  
        int mid = (begin + end) / 2;  
        SumTask leftTask = new SumTask(begin, mid, array);  
        SumTask rightTask = new SumTask(mid + 1, end, array);  
        leftTask.fork();  
        rightTask.fork();  
        return rightTask.join() + leftTask.join();  
    }  
}
```

